

Savio

An AI money companion for the one question budget apps never answer. Can I afford this, right now? It reasons over your real income, commitments and goals.



■ 27 days it says nothing ■ 3 moments it speaks and never at the checkout

Every money app answers a question nobody asked

They tell you where your salary went. That is a report on a month you cannot change any more.

The real question Can I afford this, right now? Will it break something I have already promised?

What people do instead They work it out in their head, at the moment of deciding, with no tool open.

What that costs A guess is how a month quietly goes wrong. Nobody notices until the month is over.

Ten interviews, thirty-one surveys, four countries

I went in expecting to build a smarter budget. The study said the budget was never the problem.

01

Interviews

10 long conversations about money decisions people had actually made, and regretted or not.

02

Survey

31 responses across four countries, to check whether the interview findings held at any width.

03

What I asked

Not what they track. What they do in the ninety seconds before they spend.

04

What I did not do

Test a prototype. There was nothing to test yet, and asking people to react to a screen would have taught me about the screen.

Four findings, and one of them decided the product

The strongest signal was not about tracking at all.

45% do the maths in their head

No app, no spreadsheet, nothing open. At the moment of deciding, the tool they use is memory.

71% wanted to see what a purchase does

Not what they spent last month. What this one will do to their savings. The strongest signal in the study.

Four had invented a rule


A wait-48-hours rule, a never-above-this rule. People had already built the product, badly, in their heads.

5 of 6 said decisions are slow burns

Not checkout moments. One person deliberated over a phone for a year. This is the finding I built on.

Everyone reports. Nobody answers.

Two questions separate the field. Does it know your real numbers? And does it answer the question you actually have?



One kind of user, and three jobs no app was doing



PRIMARY USER

Salaried, planning something

Earns monthly, has commitments already promised, and a goal being saved towards. Not in debt trouble. Not wealthy enough to stop thinking about it.

When I am about to spend something significant, **I want to know what it does to the rest of my month**, so I am not guessing at the worst possible moment.

When an app tells me no, **I want it to be a rule I set myself**, so I do not simply ignore it the way I ignore every other alert.

When the number is not one it can stand behind, **I want it to say so**, rather than make one up that sounds confident.

The finding scored last is the one I built on

Prioritisation scored “financial decisions are slow burns” fourth of four, 18 out of 125, because it read as a quarterly problem. I disagreed with the frequency, not the finding.

CALL	WHAT IT WAS	VERDICT
What I read	Financial decisions are slow burns, not checkout moments. Five of six interviews said so.	The finding
What the room did	Scored it last of four. 18 out of 125. A quarterly problem is not worth a product.	18/125
Where I disagreed	Not with the problem. With the frequency. The window is long, but its book-ends are monthly.	Frequency
What I ruled out	The checkout. A purchase that matters is made in ninety seconds and nobody opens an app inside them.	Cut
What I picked	Payday. The day a bonus lands. The days after a purchase, when you know whether it was worth it.	Built

The smallest thing that answers the question

The claim is that a person can decide a real purchase from one screen, without doing arithmetic. Everything that did not test that claim waited.

In scope

- + A verdict on a specific amount, with the working shown
- + Every figure computed in code, from the user's own income and commitments
- + A rule the user writes themselves, quoted back at them by name
- + Three moments to speak: payday, a windfall, a few days after a purchase
- + Reflection on purchases already made, labelled by the user
- + A refusal to answer when the numbers cannot support one

Out of scope, on purpose

- Any intervention at the checkout, ever. The first thing I fixed and the hardest to hold
- Naming a fund or a stock. That is regulated advice and needs a licence I do not have
- Tax regimes and 80C instruments. A chartered accountant's job
- Location nudging. Accurate, and it turns a companion into something that watches you
- Joint accounts. Two people's money has its own consent problems
- Real bank connections, which need Account Aggregator and an RBI-registered partner

It waits, it grounds, and it quotes you back to yourself

Three design rules, and the middle one is where the AI is allowed to work.

01 **Picks its moment** Not a popup when you have already decided. It speaks when salary lands, when a windfall arrives, and a few days after you bought something.

02 **Explains, grounded** Code produces the numbers. The model turns them into something you would act on, drawing only on what is really yours.

03 **Quotes your rule** Not “this is over budget”, but “over the ₹3,000 impulse-wait limit you set”. A rule you wrote yourself is a rule you will accept.

04 **Stays quiet when unsure** A number it cannot stand behind does not get shown. Trust built in, rather than promised in small print.

The model writes the sentence. It never counts.

The AI is not the weak half. It turns a correct number into something a person will actually act on, which is the whole difference between a calculation and a decision. Code does the arithmetic, because arithmetic has exactly one right answer.

CODE

safe to spend ₹26,532
after this ₹23,032

Works out every figure, because
arithmetic has one right answer.

MODEL

Turns a correct number into
a sentence you would act on.

IT MISCOUNTED

One verdict. The working. Your own rule.

Isn't this just a chatbot with a finance skin?

Every reviewer asks it. So the answer is a test, not a sentence. The same questions go to Savio and to a raw Gemini model, and both calls use the same underlying model, so the only variable left is the product work.

RAW GEMINI, SAME MODEL

No number

Answers from general knowledge, in one free-form blob. Will happily pick you a stock. Carries nothing between turns, and has no way to know it is inventing a number.

SAVIO

₹26,532

The real safe-to-spend, from her real state. A verdict, the trade-offs, one next step. Refuses what it is not licensed to answer, and every field is guard-verified before it is shown.


One purchase, before and after

The ninety seconds do not change. What changes is whether the decision was already made before them.

BEFORE



AFTER



What it costs, and why cost is not the constraint

A verdict is a third of a rupee. Three regulatory boundaries and one model deadline matter far more.

LINE	DETAIL	FIGURE
Per verdict	6,800 tokens in, 900 out, on Gemini 2.5 Flash.	\$0.0043
Per user	About ten model calls a month. Ten is derived from one ritual, two reflections and five considered purchases — the five is the number to argue with.	₹3.6/mo
Biggest lever	The system prompt is 61% of every input and never changes. Cached at a tenth of the rate, a verdict drops to \$0.0032.	61%
What closes the doors	Naming investments needs a SEBI licence. Tax needs a chartered accountant. Real accounts need Account Aggregator and an RBI-registered partner.	Subscription
The real cap	At ₹4 of cost, price is not the constraint. Until account connectivity lands you type your own numbers in, and a product you maintain by hand is one you stop paying for.	Data entry

Three of six calls, and what each one bought

Timing over presence

The checkout was the original idea. A purchase takes ninety seconds and nobody opens an app inside one, so I gave up the moment of maximum leverage and speak at payday instead.

Voice over vigilance

"You're near a Myntra. You've regretted four of four purchases there." Accurate, and it turns a companion into something that watches you. Gave up the sharpest nudge available.

One question over a dashboard

Stocks, funds and deposits together is a different product with different competitors. Gave up breadth to answer one thing properly.

One deadline, one boundary, one open question

Ordered by what is forced, not by what is interesting.

NEXT	WHY	WHEN
Migrate the model	Gemini 2.5 is deprecated from October 2026 and repriced after January. The model ID is already config; carrying thought signatures across six replayed turns is the work.	Jan 2027
Fix what beta found	A handful of people found what I could not. One could not work out how to leave. Another never found the Reflect tab. “Safe to spend today” sat above a whole month's figure.	5 real fixes
Real money, eventually	Nobody manages real money with it. That is a legal boundary, not an unfinished feature: Account Aggregator and RBI-registered partners, with obligations I cannot take on alone.	Blocked
The open question	Whether any of this changes what somebody does with their money. Beta testers on a demo are spending nothing. That needs users I am not allowed to have yet.	Unanswered

LIVE NOW

savio.ritikadas.in

Runs on a seeded user with six months of history. The full build write-up is at ritikadas.in/work/savio.